

The Loop AI Engine

Mock scenarios & discussion guide — Property Marketing working session

● **Attract** → ● **Engage** → ● **Convert** → ● **Optimize**

Optimize is the stage that closes the Loop — the AI engine this session is about.

Where the AI engine sits today

THE FOUR-STAGE LOOP — AS CLIENTS SEE IT ON THE PORTAL

● Attract

Paid + organic acquisition

● Engage

Property content + AEO

● Convert

Leads → tours → leases

● Optimize

AI engine + forecasts

WHAT THE OPTIMIZE ENGINE PRODUCES

30-day lease forecast

Per property, with a confidence range. Pilot property today: 10.8 leases (range 7.6–14.0).

Channel allocation & CPL

Spend, cost per lease, and projected leases for each of the five media channels.

Budget recommendations

Shift, increase, or hold — each with a rationale and an Approve & Apply / Defer choice.

AI-curated summary

Plain-language monthly digest: what changed, what's working, what needs your input.

THREE OPERATING MODES

● Auto-pilot

The engine applies small, bounded changes itself; the client reviews a weekly digest.

● Co-pilot


The engine proposes; the client or AM approves before anything is applied.

● Custom

The AM crafts a goal-driven plan; the engine executes and reports against it.

How a recommendation becomes a budget change

The pipeline behind every scenario in this deck.



Today, every budget change routes through a human — nothing goes live until a Regional Manager signs. The question for this room: where should that line sit as the engine earns trust?

Scenario 1 The budget shift

CO-PILOT

A stabilized property, \$3,400/mo across five channels. The engine spots one channel converting at 3x the cost of another.

The Ashton · 264 units · Standard tier · \$3,400/mo media

Recommendations (1)

Shift \$450/mo from Reputation → SEO

SEO is converting at ~\$200 per lease; Reputation at ~\$600. Moving 15% of Reputation spend adds a projected +2.1 leases over the next 30 days.

shift_budget

+2.1 leases (30d)

Approve & Apply

Defer

Approving routes to your team for execution — nothing changes automatically.

ASK THE ROOM

1. Would you approve this from the rationale alone — what's missing before you'd hit the button?
2. Is 15% of a channel per move the right ceiling, or should it flex by property size and season?
3. Who should see this first — the client, the AM, or both at the same moment?
4. Are there channels the engine should never be allowed to pull money from?

Guardrails in the build today: a move is capped at 15% of the source channel, and the engine stays silent (“the Loop is balanced”) when the shift is under \$50 or the projected gain is under half a lease.

Scenario 2 Asking for more money

CO-PILOT

Google Ads is capped by budget while performance status is yellow. The engine wants to ask the client for more.

Recommended budget change

STATUS: YELLOW

Raise Paid Search \$1,200 → \$1,800/mo

Paid Search is losing ~38% of available impressions to budget while marketing status is yellow. Raising spend recovers most of that lost share (capped at +50% per step).

budget_change

+\$600/mo

Available impressions

62% captured

38% lost to budget

Approving builds the deal + quote for your Regional Manager to sign — nothing goes live until signed.

Guardrails in the build today: fires only when performance status is red or yellow and $\geq 10\%$ of impression share is lost to budget; capped at +50% and \$10k/mo; suppressed when a budget deal is already in flight.

ASK THE ROOM

1. What evidence earns the right to ask a client for more budget — and is lost impression share enough?
2. Is +50% in one step too aggressive? Should increases ladder in smaller moves?
3. When should the engine explicitly not ask, even when the data says yes?
4. Does this ask belong on the portal, or in the AM's hands first?

Scenario 3 Auto-pilot acts alone

AUTO-PILOT

A Local-tier property with a small budget. The engine makes a bounded move itself and tells the client afterward.

Your weekly summary

AI-CURATED

What we're handling

"The Loop adjusted spend by \$300 toward Paid Search this week after paid social cost-per-lease rose for the third straight month. No action needed from you — full detail is in your Loop timeline."

≤ 15%

of the source
channel

≤ \$500

per automatic move

+ forecast

impact must be
positive

7 days

warm-up on new
properties

Auto-pilot can only self-approve budget shifts, SEO refreshes, and AEO content batches — never budget increases.

ASK THE ROOM

1. Which properties should default to auto-pilot — by tier, budget size, or client relationship?
2. Is a \$500 cap per automatic move too timid, or too bold?
3. How loudly do we tell the client the AI acted — weekly digest, same-day email, or only on request?
4. What single action, if automated, would save your team the most hours?

Proposed default in ADR 0009: Local & Lite tiers start in auto-pilot, Basic & Standard in co-pilot, Premium in custom. This room gets a vote.

Scenario 4 The forecast misses

ACCOUNTABILITY

Last month the engine forecast 10.8 leases with an 80% confidence range of 7.6–14.0. The property signed 5.

Forecast vs. actual — last 30 days

10.8

forecast leases

5

actual leases signed

OUTSIDE RANGE

80% confidence range: 7.6 – 14.0



The engine already scores itself: average miss size and how often actuals land inside the range, per property and portfolio-wide. An AM alert on big deviations is designed but not yet switched on — this room sets the threshold.

ASK THE ROOM

1. At what size of miss should an AM be alerted — before the client notices?
2. Should clients see forecast accuracy? Is it confidence-building or ammunition?
3. What should next month's AI summary say about the miss?
4. What offline factors — renovations, staffing, pricing moves — should the engine be told about, and by whom?

Scenario 5 When more spend won't fix it

JUDGMENT CALL

Lead volume is healthy; leases aren't following. The engine's read: the problem isn't marketing.

AI assessment — external factors

THREAT

“The property is priced ~8% above its comp set with no concessions; two nearby comps are offering six weeks free. Additional media spend is unlikely to convert at acceptable cost until pricing or concessions move.”

142

leads (30d)

31

tours

4

leases

Portal callout shown with it: “Your property's health score suggests operational improvements may be needed alongside marketing investment to fully capture lead volume.”

ASK THE ROOM

1. Should the AI ever tell a client “more budget won't help”? How direct can it be?
2. Who does the engine hand this to — operations, pricing, or the AM to deliver in person?
3. What non-marketing data (pricing, comps, concessions) would you trust it to reason about?
4. How do we keep this from reading as excuse-making when performance is soft?

This mirrors the real pattern the spec names: high lead volume, low lease conversion. Today the AI flags it in the internal SWOT; the open question is what the client-facing version says.

Scenario 6 “30 leases by October 1”

CUSTOM MODE

A lease-up with a hard deadline. Custom mode works backwards from the goal — and finds the funnel runs dry.

Full-funnel plan — working backwards from the goal

30 leases needed

~**960** qualified leads required

~**48k** site sessions required

~**2.4M** impressions required

Where it runs dry: Paid Search saturates at ~1.2× current spend — past that cap, extra dollars are only ~12% effective. The plan pushes incremental budget into Meta, CTV, and retargeting, which still have headroom.

ASK THE ROOM

1. Do goals like this actually reach the marketing team today — and early enough to act?
2. Who owns the plan: the engine drafts and the AM edits, or the reverse?
3. When the goal is unreachable at any budget, what should the engine say — and to whom?
4. Which upper-funnel channels do you trust when the deadline is real?

This is the Budget Simulator already on the portal (by channel · by category · scenarios). Custom mode wires it to a named goal with a deadline — the engine reports weekly against the goal, not against last month.

The bigger questions

Beyond individual scenarios — the answers here shape the engine's defaults for all 700+ properties.

● Guardrails & trust

- Are the current caps right — 15% per shift, \$500 auto-pilot, +50% per increase?
- What should the engine never do without a named human owning it?
- When data is thin (new property, <3 months history), should it stay silent or hedge loudly?

● Client communication

- How do we label AI-authored content — and is “AI-curated” a feature or a liability?
- What tone earns trust: confident, hedged, or numbers-only?
- What should a client never learn from the portal before their AM does?

● Approvals & workflow

- Real-time approvals or a weekly batch queue for AMs?
- Should AMs be able to counter-propose (edit the number) rather than just approve/reject?
- Whose approval carries the day when client and AM disagree?

● Field reality

- What are the 3 situations that burned the most of your time last quarter?
- Which recurring client questions could the engine answer before they're asked?
- What does the engine get wrong about how budgets actually move at RPM?

What we need from this room

1

Pressure-test the guardrails

The caps in this deck — 15% per shift, \$500 auto-pilot moves, +50% budget increases — are engineering defaults, not marketing judgment. Tell us where they're wrong.

2

Give us three real situations

The messiest budget calls, client escalations, or missed opportunities from last quarter. We'll replay them through the engine and show you what it would have done.

3

Nominate pilot properties

5–10 properties for co-pilot mode — enough variety in tier, market, and budget size to stress the recommendations before a broader rollout.

The forecasting, recommendation, and approval flows in this deck are built and running on pilot data today. Your answers set the guardrails before we switch them on broadly.